

Deep Analysis of Clinical Summarization Papers Listed in the Uploaded Document

Scope and source limitations

Your uploaded document lists four published studies and one proposed MSK study. I analyzed the four papers and did **not** treat the “Proposed MSK Study” as a published paper because the document presents it as a future study design rather than an existing manuscript. I was able to independently verify and inspect public article records for **Van Veen et al.** and **Sivarajkumar et al.** For **Navarro et al.** and **Bednarczyk et al.**, I could not confidently match the exact public manuscripts from the title/author strings supplied in the document, so those two summaries are grounded in the uploaded document’s internal study notes rather than a direct paper read. [filecite turn0file0](#)

Navarro et al

Paper Overview

The uploaded document identifies this paper as “**Expert evaluation of LLMs for clinical dialogue summarization**” by **Navarro et al.**, with publication venue listed as **Scientific Reports (2025)**. No institution is visible in the uploaded source summary. [filecite turn0file0](#)

Abstract Summary

Based on the uploaded summary, this paper appears to ask a practical question: when clinicians judge summaries of clinical dialogue, do current LLMs produce notes that are actually useful, and do common automatic metrics such as ROUGE reflect that usefulness? The central message is that at least in this dataset, clinician judgment favored ChatGPT-style summaries more than ROUGE did, suggesting that automatic overlap metrics may miss clinically meaningful quality. [filecite turn0file0](#)

Key Contributions

- It centers **clinician evaluation** rather than relying only on automated text-overlap metrics for judging summary quality. [filecite turn0file0](#)
- It highlights an apparent **disconnect between ROUGE and clinician preference**, with ChatGPT reportedly performing best in human review but poorly on ROUGE. [filecite turn0file0](#)
- It argues that future evaluation should explicitly track **omissions** and **hallucinations**, not just generic summary similarity. [filecite turn0file0](#)

Methodology

- **Study design:** comparative evaluation of LLM-generated clinical dialogue summaries against human expectations in a clinical setting. [\[1\]](#)
- **Data collection methods:** the uploaded summary reports a dataset of **27 dialogues**, with **5 dialogues evaluated**, and **5 clinicians** rating the summaries. [\[2\]](#)
- **Analysis techniques:** clinicians scored summaries on **coherence, fluency, and usefulness**, and these ratings were considered alongside **ROUGE** scores. [\[3\]](#)
- **Tools/frameworks used:** the available source summary does not specify the exact LLM versions, prompt templates, or software stack. [\[4\]](#)

Key Findings

- **ChatGPT performed best in clinician evaluation** on coherence, fluency, and usefulness. [\[5\]](#)
- The same model reportedly performed **poorly on ROUGE**, suggesting that ROUGE may not be a reliable stand-alone metric for clinical dialogue summarization. [\[6\]](#)
- The study concluded that **LLM summaries approached human quality**, at least in this small primary care sample. [\[7\]](#)

Figures and Tables Analysis

No figures or tables from the full manuscript are visible in the material I could access. From the summary alone, the most likely key visuals would be a **table of clinician ratings by model** and a **comparison table between ROUGE and human judgment**. Because those pages were not available, I cannot verify exact axes, confidence intervals, or model-by-model score distributions. [\[8\]](#)

Technical Concepts Explained

ROUGE is a family of summary metrics based on overlap between a model's output and a reference summary. In practice, ROUGE rewards textual similarity, which can be useful but may fail when two summaries use different wording to express the same clinically correct point. That is exactly the kind of mismatch this paper seems to be flagging. [\[9\]](#)

In this context, an **omission** means leaving out clinically important information, while a **hallucination** means introducing unsupported or false medical content. Those are much closer to patient-safety concerns than surface-level wording overlap. [\[10\]](#)

Strengths

The paper's biggest strength, as presented in your document, is that it asks the right evaluation question for a clinical workflow: **which summary would a clinician actually trust and use?** That is more clinically meaningful than asking only whether a model matched a reference sentence token-for-token. [\[11\]](#)

Limitations

The limitations are substantial and clearly acknowledged in the uploaded summary: **small sample size**, **primary care context**, uncertain generalizability to other specialties such as oncology, and **no safety outcomes** linked to downstream clinical decisions. That means the paper is hypothesis-generating rather than practice-changing. [filecite turn0file0](#)

Practical Implications

For real-world deployment, this paper supports using **expert human review as a primary evaluation layer** and treating ROUGE as, at most, a supportive metric. It also argues for auditing summaries for omission and hallucination before introducing them into a live clinical workflow. [filecite turn0file0](#)

TL;DR

This study suggests that LLMs can produce clinical dialogue summaries that clinicians often judge as high quality, but standard automatic metrics may underrate them. The practical takeaway is that clinical summarization should be judged by **clinical usefulness and safety**, not by ROUGE alone. [filecite turn0file0](#)

Bednarczyk et al

Paper Overview

The uploaded document identifies this paper as “**Scientific Evidence for Clinical Text Summarization Using LLMs: Scoping Review**” by **Bednarczyk et al.**, with venue listed as **JMIR (2025)**. No institution is visible in the uploaded summary. [filecite turn0file0](#)

Abstract Summary

This paper is presented as a **scoping review** of the evidence base for LLM-based clinical text summarization. Its plain-English conclusion is that the field is still early: most studies are retrospective, clustered around a few benchmark settings such as radiology, ICU, and MIMIC-derived data, and rarely include the kinds of external validation, safety assessment, or bias analysis needed for clinical deployment. [filecite turn0file0](#)

Key Contributions

- It maps the literature and shows that the evidence base for LLM clinical summarization is still **immature and heavily retrospective**. [filecite turn0file0](#)
- It identifies topic concentration around **radiology, ICU settings**, and **MIMIC-based datasets**, implying narrow coverage of real clinical use cases. [filecite turn0file0](#)
- It makes the field’s biggest evaluation gaps explicit: **very little external validation**, almost **no safety analysis**, and **no bias analyses** in the reviewed set. [filecite turn0file0](#)

Methodology

- **Study design:** scoping review of published LLM-based clinical text summarization studies. [filecite turn0file0](#)
- **Data collection methods:** the source summary states that **30 studies** were reviewed, but it does not expose the search databases, screening flow, or inclusion criteria. [filecite turn0file0](#)
- **Analysis techniques:** descriptive synthesis of the literature, with emphasis on validation setting, safety work, and evidence gaps. [filecite turn0file0](#)
- **Tools/frameworks used:** because the full article was not available to inspect, I cannot verify whether the review followed PRISMA-ScR or another formal review framework. [filecite turn0file0](#)

Key Findings

- Of the **30 studies reviewed**, only **2** included **external validation**. [filecite turn0file0](#)
- Only **1** study included a **safety analysis**. [filecite turn0file0](#)
- **0** studies included a **bias analysis**, which is striking for a clinical AI literature. [filecite turn0file0](#)

Figures and Tables Analysis

No figures or tables from the full review were visible in the sources I could inspect. In a complete scoping review, I would normally expect a **study-selection flow diagram**, a **characteristics-of-included-studies table**, and a **gap map** showing where evidence is concentrated versus missing. Because those pages were not accessible, I cannot comment on exact diagram counts beyond what is stated in the uploaded summary. [filecite turn0file0](#)

Technical Concepts Explained

A **scoping review** is meant to map what research exists, how it is distributed, and where major gaps remain. It usually answers “what has been studied?” more than “what definitely works?” That makes it a good design for a young field like clinical LLM summarization. [filecite turn0file0](#)

External validation means testing a method on data from a different setting than the one used to build it. In healthcare, that matters because performance often drops when a model moves from one institution, workflow, or patient population to another. A **bias analysis** checks whether errors or performance vary unfairly across groups or contexts. [filecite turn0file0](#)

Strengths

This paper’s strength is that it zooms out from single-model hype and asks a more durable question: **how convincing is the actual evidence base?** The answer, as summarized in your document, is sobering and useful. [filecite turn0file0](#)

Limitations

Because I could not inspect the full article, I cannot verify the databases searched, the date ranges, or the exact screening strategy. So my analysis here is constrained to the uploaded summary rather than the review manuscript itself. [filecite turn0file0](#)

Practical Implications

For anyone planning a specialty-specific study, this paper is highly actionable. It implies that the next valuable papers should emphasize **prospective evaluation, external validation, safety endpoints, bias auditing, and workflow outcomes** rather than yet another retrospective benchmark on MIMIC alone. [filecite turn0file0](#)

TL;DR

This scoping review says the field of LLM-based clinical summarization is promising but scientifically underdeveloped. Most studies are retrospective and narrow, and the literature still lacks the validation, safety, and fairness analyses needed for clinical trust. [filecite turn0file0](#)

Van Veen et al

Paper Overview

The public record identifies this paper as **“Adapted Large Language Models Can Outperform Medical Experts in Clinical Text Summarization”** by Dave Van Veen, Cara Van Uden, Louis Blankemeier, Jean-Benoit Delbrouck, Asad Aali, Christian Bluethgen, Anuj Pareek, Malgorzata Polacin, Eduardo Pontes Reis, Anna Seehofnerová, Nidhi Rohatgi, Poonam Hosamani, William Collins, Neera Ahuja, Curtis P. Langlotz, Jason Hom, Sergios Gatidis, John Pauly, and Akshay S. Chaudhari. The journal venue is **Nature Medicine**, published **27 February 2024**; the author affiliations include multiple Stanford units, the Stanford Center for Artificial Intelligence in Medicine and Imaging, UT Austin, University Hospital Zurich, Copenhagen University Hospital, and Albert Einstein Israelite Hospital in São Paulo. The uploaded document’s “2023” label matches the paper’s arXiv/preprint timing rather than the journal publication date.

2

Abstract Summary

This is the strongest and most directly substantiated paper in your set. It evaluates how well adapted LLMs summarize several kinds of clinical text, then goes beyond benchmark metrics by asking physicians which summaries they actually prefer; in that controlled setting, the best adapted model was often judged as comparable to or better than medical experts. [3](#)

Key Contributions

- It benchmarks **eight LLMs** across **four summarization tasks** using **six datasets**, rather than focusing on a single narrow task. [4](#)
- It includes a **clinical reader study with 10 physicians**, which directly evaluates completeness, correctness, and conciseness. [5](#)

- It adds a **safety analysis tied to potential medical harm** and compares quantitative metrics with physician preferences, showing that the usual NLP metrics only weakly track clinical judgment. ⁶

Methodology

- **Study design:** multi-stage evaluation pipeline. The authors first benchmarked model–adaptation combinations quantitatively, then ran a blinded physician A/B reader study comparing the best model against medical expert summaries, and finally performed a safety analysis of clinically meaningful errors. The study framework is shown in **Figure 1**. ⁷
- **Data collection methods:** they used six public datasets spanning four tasks—**Open-i**, **MIMIC-CXR**, and **MIMIC-III** for radiology reports; **MeQSum** for patient questions; **ProbSum** for progress notes; and **ACI-Bench** for doctor–patient dialogue. Reported dataset sizes ranged from **126 dialogue samples** used from ACI-Bench to **128K radiology reports** in MIMIC-CXR. ⁸
- **Analysis techniques:** for quantitative evaluation, the paper used **BLEU**, **ROUGE-L**, **BERTScore**, and **MEDCON**. Human evaluation used a **five-point Likert scale** for completeness, correctness, and conciseness, with **Wilcoxon signed-rank testing** and reader reliability estimates. Safety analysis asked physicians to rate the likely **extent** and **likelihood** of harm if the worse summary were used clinically. ⁹
- **Tools/frameworks used:** the workflow used **in-context learning** with nearest-neighbor example selection via **PubMedBERT**, **QLoRA** for parameter-efficient fine-tuning, **PyTorch**, **QuickUMLS** for MEDCON, and consumer GPU hardware including an **NVIDIA Quadro RTX 8000**. ¹⁰

Key Findings

- The authors identified **GPT-4 with the maximum allowable number of in-context examples** as their best overall model/method combination. ¹¹
- In the reader study, medical expert summaries were preferred in only **19%** of cases, while model summaries were **non-inferior in 45%** and **preferred in 36%**. ¹²
- In the harm analysis, the worse medical-expert summary was associated with a higher estimated **likelihood of harm (14%)** and **extent of harm (22%)** than the worse model summary (**12%** and **16%**, respectively). ¹³
- Among open-source models, **FLAN-T5** emerged as the strongest QLoRA baseline, and **medical Q&A tuning did not automatically help summarization**: Med-Alpaca underperformed Alpaca on this task family. ¹⁴
- With enough in-context examples, proprietary models—especially **GPT-3.5** and **GPT-4**—substantially outperformed the other model classes; by contrast, simple model size alone was not a guarantee of better summarization. ¹⁵
- **BERTScore** and **MEDCON** correlated best with correctness, but the correlations were still low, around **0.2**, reinforcing that automated metrics alone are not enough for clinical-readiness claims. ¹⁶

Figures and Tables Analysis

- **Figure 1: Framework overview.** This figure visually lays out the paper’s whole logic: model/task benchmarking first, clinician reader study second, safety analysis third. That matters because it shows the paper is not just a leaderboard exercise; it is explicitly trying to bridge model performance with clinical acceptability and risk. ¹⁷
- **Figure 2: Model prompts and temperature.** The visual combines prompt anatomy with a small results table and shows a practical finding: **lower temperature** and assigning the model **clinical**

expertise improved summary performance relative to less relevant expertise framing. This supports the paper’s claim that prompt design is an adaptation mechanism, not a cosmetic detail. ¹⁸

- **Figure 3: Identifying the best model/method.** This is a composite figure. It shows Med-Alpaca underperforming Alpaca on summarization, compares **ICL versus QLoRA** on open-source models, shows performance improving as the number of in-context examples grows, and ends with head-to-head win-rate heatmaps where **GPT-4** generally dominates. The figure strongly supports the claim that **task adaptation** and **context length** matter more than generic “medical” branding. ¹⁹
- **Figure 4: Clinical reader study.** This is the paper’s most consequential figure. It shows pooled physician preference distributions and harm ratings, visually reinforcing that the best model scored higher on completeness, correctness, and conciseness, and that the harm profile favored the model over the human comparator in this experiment. ²⁰
- **Figure 5: Annotation examples for radiology reports.** This figure presents qualitative side-by-side exemplars with highlighted omissions and incorrect concepts. It is important because it makes the safety problem concrete: not all errors are equally dangerous, and the clinically relevant distinction is what was omitted or misstated, not just whether a summary “looks fluent.” ²¹
- **Extended Data Table on datasets and task instructions.** The table gives the paper’s task breadth at a glance: radiology, questions, notes, and dialogue vary enormously in token length and lexical diversity, and the prompt instructions differ by task. That directly supports the authors’ argument that clinical summarization is really a **family of tasks**, not one uniform benchmark. ²²

Technical Concepts Explained

In-context learning means giving a model a few task-relevant examples inside the prompt instead of changing its weights. **QLoRA** is a memory-efficient way to fine-tune only a small subset of parameters rather than retraining the full model. In high-level terms, ICL adapts at inference time; QLoRA adapts by lightweight training. ²³

The paper uses four evaluation lenses: **BLEU** and **ROUGE-L** for surface-level overlap, **BERTScore** for semantic similarity, and **MEDCON** for agreement in extracted medical concepts using **QuickUMLS**. The fact that these correlated only weakly with physician preference is one of the paper’s biggest methodological lessons. ²⁴

Strengths

This is a very strong paper methodologically because it has breadth of tasks, breadth of models, human expert evaluation, and an explicit safety frame. It also resists a simplistic “bigger model wins” narrative by showing nuanced trade-offs among model family, adaptation strategy, prompt design, and context length. ²⁵

Limitations

Even this paper remains a controlled benchmark-plus-reader-study design rather than a prospective clinical workflow trial. The **dialogue task was excluded from the reader study** because it was too cumbersome for readers, and the best-performing system depends on **GPT-4**, a proprietary model whose deployment characteristics may differ across institutions and over time. ²⁶

Practical Implications

For deployment, the paper supports a model of use in which LLMs generate **candidate summaries for clinician review**, especially when long context and relevant in-context examples are available. It also argues that hospitals should monitor not just generic summary quality but **clinical completeness, fabrication risk, and correction burden**, because those are the dimensions that matter most once models become “good enough” to look convincing. ²⁷

TL;DR

Van Veen et al. is a major paper in this space: it shows that adapted LLMs—especially GPT-4 with strong in-context examples—can produce clinical summaries that physicians often judge as at least as good as expert-written ones in controlled comparisons. But the deeper lesson is not “LLMs solved summarization”; it is that **adaptation helps, human evaluation is essential, and safety must be measured explicitly**. ²⁸

Sivarajkumar et al

Paper Overview

The accessible public record identifies the paper as “**An Empirical Evaluation of Prompting Strategies for Large Language Models in Zero-Shot Clinical Natural Language Processing**” by **Sonish Sivarajkumar, Mark Kelley, Alyssa Samolyk-Mazzanti, Shyam Visweswaran, and Yanshan Wang**, posted to arXiv on **14 September 2023**. The uploaded document shortens the title to “**An Empirical Evaluation of Prompting Strategies for Large Language Models in Clinical Tasks**” and labels it as a **2024** paper, which likely reflects either a shortened citation form or later manuscript circulation rather than the initial public preprint date. An institutional affiliation was not visible on the accessible arXiv abstract page. ²⁹

Abstract Summary

This paper is not narrowly about summarization; it is a broader prompt-engineering study across zero-shot clinical NLP tasks. In plain English, it asks a question that matters a lot for deployment: if you keep the underlying LLM fixed but change the prompt strategy, how much does clinical-task performance move—and what kinds of prompts are most reliable? ³⁰

Key Contributions

- It presents one of the early **systematic comparisons of prompt engineering strategies** for clinical NLP in the generative-AI era. ³⁰
- It evaluates prompting across **five clinical NLP tasks** and **three LLMs** rather than reporting a single anecdotal prompt win. ³¹
- It tests prompt types from prior work and adds **heuristic prompting** and **ensemble prompting**, while also comparing **zero-shot** with **few-shot** prompting. ³¹

Methodology

- **Study design:** comparative experimental benchmark of prompting strategies for fixed LLMs on zero-shot clinical NLP tasks. ³¹
- **Data collection methods:** the abstract identifies five task families—**clinical sense disambiguation, biomedical evidence extraction, coreference resolution, medication status extraction**, and

medication attribute extraction—but it does not expose the exact datasets or sample counts on the accessible page. ³¹

- **Analysis techniques:** the public abstract confirms comparison across prompt styles and across zero-shot versus few-shot setups, but it does not show the exact task metrics or full result tables in the accessible record I could inspect. ³¹
- **Tools/frameworks used:** the tested models were **GPT-3.5, BARD, and LLAMA2**; the prompt families included **simple prefix, simple cloze, chain-of-thought, anticipatory prompts, heuristic prompting, and ensemble prompting.** ³¹

Key Findings

- The uploaded document summarizes the paper’s main result as: **prompting strategy significantly impacts performance, and structured prompting improves reliability.** ³⁰
- The same document states that when the authors compared **zero-shot, few-shot, and structured prompting, the structured methods performed best.** ³⁰
- The uploaded summary also notes that **errors persisted despite optimization**, meaning better prompts reduced but did not eliminate clinically important failure modes. ³⁰
- The public abstract supports the broader conclusion that the paper offers **guidelines for prompt engineering in clinical NLP**, but exact statistics were not visible on the accessible page. ³⁰

Figures and Tables Analysis

I could not directly inspect figures or tables for this paper from the accessible public sources. Based on the abstract and the uploaded summary, the full paper likely includes **per-task performance tables, prompt-family ablations, and possibly model-by-prompt interaction charts.** Because those pages were not accessible, I cannot verify exact effect sizes, confidence intervals, or the specific visual evidence behind the “structured prompting is best” claim. ²⁹

Technical Concepts Explained

A **zero-shot** prompt asks the model to perform a task without examples. A **few-shot** prompt includes a small number of labeled examples in the prompt to teach the pattern. A **chain-of-thought** prompt asks the model to reason step by step before answering, while an **ensemble prompt** usually combines outputs or reasoning paths to improve robustness. A **heuristic prompt** incorporates manually designed guidance or constraints that reflect domain knowledge. ³¹

The practical message is simple: with LLMs, the **prompt is part of the model behavior.** Two prompts given to the same base model can act like meaningfully different systems, which is why prompt design matters for safety and reproducibility. ²⁹

Strengths

The paper’s strength is its framing. Rather than treating prompt writing as an afterthought, it studies prompting as an experimental variable with clinical consequences. It also spans multiple task types, making the findings more useful than a one-task prompt anecdote. ³⁰

Limitations

The accessible public record did not expose the full results tables, which limits how much quantitative detail I can safely report. More importantly, the uploaded summary notes that **hallucinations and omissions still occurred** and that there was **limited real-world validation**, so the work should be read as guidance for controlled benchmarking, not as proof of deployment readiness. [filecite turn0file0](#)

Practical Implications

For clinical AI teams, this paper supports treating prompt templates like versioned software components: they should be **designed, benchmarked, documented, and regression-tested**. It also supports structured or multi-step prompting as a plausible low-cost safety control before moving to heavier model retraining.

TL;DR

Sivarajkumar et al. shows that prompt engineering is not cosmetic in clinical NLP—it is performance-critical. The paper’s practical takeaway is that **structured prompting can make LLM behavior more reliable**, but prompt optimization alone does not eliminate clinical errors. [29](#)

Cross-paper synthesis

Taken together, these papers tell a coherent story. **LLMs can be genuinely strong clinical summarizers in controlled settings**, especially when adapted with task-aware prompts or in-context examples, but **raw benchmark scores are not enough to establish clinical readiness**. Both the Navarro summary and the Van Veen paper point toward the same lesson: clinician judgment and safety-oriented error analysis matter more than ROUGE alone. [32](#)

A second recurring theme is that **adaptation matters**. Van Veen et al. shows strong gains from in-context learning and careful task adaptation, while Sivarajkumar et al. argues that prompt structure itself materially changes performance and reliability. So across the set, the best interpretation is that clinical summarization quality is driven less by “which LLM?” in isolation and more by the **whole pipeline**: prompt design, context selection, model choice, evaluation method, and human oversight. [33](#)

A third theme is that the **evidence base remains thin where it matters most**. The Bednarczyk review summary says the literature is still mostly retrospective and narrow, with very little external validation, almost no safety work, and no bias analyses in the reviewed set. That makes it hard to generalize from benchmark wins to safe deployment in specialty workflows such as surgical breast oncology.

[filecite turn0file0](#)

For your domain specifically, the “Proposed MSK Study” in the uploaded document is well aligned with the gaps exposed by the literature: prospective comparison, specialty-specific context, explicit taxonomy of **omission, hallucination, misinterpretation, contradiction, and temporal errors**, plus endpoints such as **clinically significant error rate, correction burden, trust, and usability**. As a research direction, that is exactly the kind of study these papers collectively imply is now needed. [filecite turn0file0](#)

2 3 5 27 28 32 <https://doi.org/10.1038/s41591-024-02855-5>

<https://doi.org/10.1038/s41591-024-02855-5>

7 17 [https://media.springernature.com/m312/springer-static/image/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig1_HTML.png)

[art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig1_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig1_HTML.png)

[https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig1_HTML.png)

[41591_2024_2855_Fig1_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig1_HTML.png)

18 [https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig2_HTML.png)

[MediaObjects/41591_2024_2855_Fig2_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig2_HTML.png)

[https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig2_HTML.png)

[41591_2024_2855_Fig2_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig2_HTML.png)

19 [https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig3_HTML.png)

[MediaObjects/41591_2024_2855_Fig3_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig3_HTML.png)

[https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig3_HTML.png)

[41591_2024_2855_Fig3_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig3_HTML.png)

20 [https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig4_HTML.png)

[MediaObjects/41591_2024_2855_Fig4_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig4_HTML.png)

[https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig4_HTML.png)

[41591_2024_2855_Fig4_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig4_HTML.png)

21 [https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig5_HTML.png)

[MediaObjects/41591_2024_2855_Fig5_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig5_HTML.png)

[https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig5_HTML.png)

[41591_2024_2855_Fig5_HTML.png](https://media.springernature.com/m312/springer-static/image/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Fig5_HTML.png)

22 [https://media.springernature.com/lw867/springer-static/esm/art%3A10.1038%2Fs41591-024-02855-5/](https://media.springernature.com/lw867/springer-static/esm/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Tab2_ESM.jpg)

[MediaObjects/41591_2024_2855_Tab2_ESM.jpg](https://media.springernature.com/lw867/springer-static/esm/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Tab2_ESM.jpg)

[https://media.springernature.com/lw867/springer-static/esm/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/](https://media.springernature.com/lw867/springer-static/esm/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Tab2_ESM.jpg)

[41591_2024_2855_Tab2_ESM.jpg](https://media.springernature.com/lw867/springer-static/esm/art%3A10.1038%2Fs41591-024-02855-5/MediaObjects/41591_2024_2855_Tab2_ESM.jpg)

29 30 31 <https://arxiv.org/abs/2309.08008>

<https://arxiv.org/abs/2309.08008>