



Why Scaling AI Evaluation is Harder Than You Think (and How to Fix It)

Based on 1 source



Why Scaling AI Evaluation is Harder Than You Think (and How to Fix It)

1. Introduction: The Scalability Wall

If you've built an LLM-powered application, you've hit the "vibe check" wall. It's that moment where your demo looks great, but as soon as you move toward production, you realize you have no idea if your latest prompt tweak actually improved things or just broke three other edge cases.

Evaluating LLM outputs at scale is one of the hardest problems in modern AI engineering. Human evaluation is the gold standard, but it doesn't scale to thousands of requests.



Google NotebookLM

Install App

INSTALL



quality; they care about word overlap, whereas your users care about whether the answer was actually helpful or safe.

The solution is the **LLM-as-a-Judge** pattern: using a high-capability model (like GPT-4o) to evaluate the output of your application's model. In this guide, I'll break down how to move past

Generated by AI based on user-provided sources and instructions

[Terms of Service](#) | [Privacy Policy](#) | [Learn more](#)