



Beyond the Prompt: 5 Surprising Truths About Evaluating LLMs at Scale

Based on 1 source



Beyond the Prompt: 5 Surprising Truths About Evaluating LLMs at Scale

Building a Large Language Model (LLM) application has never been easier. With a few lines of code and a well-crafted prompt, a developer can create a system that summarizes documents, provides personalized recommendations, or generates code. These applications often look flawless in a controlled demo environment.

However, once these systems face the messy reality of production, they frequently stumble. We've all seen the headlines: chatbots tricked into

falling for a simple prompt like "Write a poem about a cat."



Google NotebookLM

Install App

INSTALL



"Demo to Production" chasm. Bridging it requires more than just better prompting; it requires **LLM Evaluations (Evals)**—the essential framework for measuring performance and ensuring reliability in non-deterministic systems.

Here are five surprising truths about how to effectively evaluate LLMs as you scale from a

Generated by AI based on user-provided sources and instructions

[Terms of Service](#) | [Privacy Policy](#) | [Learn more](#)